

ACED Simple Regression

```
library(DescTools)
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.2      v readr      2.1.4
v forcats    1.0.0      v stringr    1.5.0
v ggplot2    3.4.3      v tibble     3.2.1
v lubridate  1.9.2      v tidyr      1.3.0
v purrr      1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

ACED Data

```
ACEDextract <- read_csv("ACED_extract1.csv",na="-999")
```

```
Rows: 290 Columns: 29
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr  (7): SubjID, Session, Cond_code, Sequencing, Feedback, Gender, Level_Code
```

```
dbl (22): Correct, Incorrect, Reamaining, ElapsedTime, Race, pre_scaled, pos...
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```

ACEDextract$Session <- factor(ACEDextract$Session)
ACEDextract$Cond_code <- factor(ACEDextract$Cond_code)
ACEDextract$Sequencing <- factor(ACEDextract$Sequencing)
ACEDextract$Feedback <- factor(ACEDextract$Feedback)
ACEDextract$Gender <- factor(ACEDextract$Gender)
ACEDextract$Race <- factor(ACEDextract$Race,1:8)
ACEDextract$Level_Code <- factor(ACEDextract$Level_Code)

```

```

ACEDextract %>%
  mutate(gain=post_scaled-pre_scaled) ->
  ACEDextract

```

Research Questions

In this case study we will address the first research question.

1. Do the pretest, posttest and internal game measures measure the same thing? (Validity and Reliability)

Making Scatterplots

Use `geom_point()` with `ggplot()` to make a scatterplot.

Scatterplot


Here is a simple scatterplot.

```

EAPxCorrect <- ggplot(ACEDextract,aes(x=Correct,y=EAP.sgp)) +
  geom_point()
EAPxCorrect

```

Warning: Removed 60 rows containing missing values (``geom_point()``).



Adding lines and smooths

The function `geom_smooth()` adds a smooth line.

A few key arguments:


- `method` – “lm”, “loess”, “glm”, “gam”
- `formula` – This allows specifying other kinds of curves.
- `na.rm` – Logical, if TRUE then suppresses warning about NAs
- `se` – Logical, default TRUE, should standard errors be plotted.

```
EAPxCorrect + geom_smooth(method="lm", se=FALSE)
```

```
`geom_smooth()` using formula = 'y ~ x'
```

```
Warning: Removed 60 rows containing non-finite values (`stat_smooth()`).
```


```
Warning: Removed 60 rows containing missing values (`geom_point()`).
```



```
EAPxCorrect + geom_smooth(method="lm")
```

```
`geom_smooth()` using formula = 'y ~ x'
```


```
Warning: Removed 60 rows containing non-finite values (`stat_smooth()`).  
Removed 60 rows containing missing values (`geom_point()`).
```



```
EAPxCorrect + geom_smooth(method="loess")
```

```
`geom_smooth()` using formula = 'y ~ x'
```

```
Warning: Removed 60 rows containing non-finite values (`stat_smooth()`).  
Removed 60 rows containing missing values (`geom_point()`).
```




Jittering

When the data are integers (as in the count), sometimes points plot on top of each other.

Jittering (adding a bit of random noise) can help.

```
ggplot(ACEDextract, aes(x=Correct, EAP.sgp)) +  
  geom_point(position="jitter")
```

Warning: Removed 60 rows containing missing values (`geom_point()`).



Coloring points


Attach a factor variable to

- `color` (line color) or `fill` (interior color)
- `shape` shape of plotting symbol
- `linetype` – type of the line (solid, dotted, dashed, &c).

Note: color can be a problem if (a) printing graph in black and white, or (b) show to somebody with limited color perception (about 8% of the population). Try to pair color with another aesthetic (e.g., shape or linetype).

```
ggplot(ACEDextract, aes(x=Correct, EAP.sgp, color=Level_Code,
                        shape=Level_Code)) +
  geom_point()
```

Warning: Removed 60 rows containing missing values (``geom_point()``).



Calculating a correlation

Var and Cor

```
var(ACEDextract$Correct,ACEDextract$EAP.sgp)
```

```
[1] NA
```

```
cor(ACEDextract$Correct,ACEDextract$EAP.sgp)
```

```
[1] NA
```

```
select(ACEDextract,starts_with("EAP")) %>% cor()
```

	EAP.sgp	EAP.cr	EAP.dt	EAP.eg	EAP.exp	EAP.ext	EAP.mod	EAP.rr	EAP.tab
EAP.sgp	1	NA	NA	NA	NA	NA	NA	NA	NA
EAP.cr	NA	1	NA	NA	NA	NA	NA	NA	NA

EAP.dt	NA	NA	1	NA	NA	NA	NA	NA	NA
EAP.eg	NA	NA	NA	1	NA	NA	NA	NA	NA
EAP.exp	NA	NA	NA	NA	1	NA	NA	NA	NA
EAP.ext	NA	NA	NA	NA	NA	1	NA	NA	NA
EAP.mod	NA	NA	NA	NA	NA	NA	1	NA	NA
EAP.rr	NA	NA	NA	NA	NA	NA	NA	1	NA
EAP.tab	NA	NA	NA	NA	NA	NA	NA	NA	1
EAP.vr	NA	NA	NA	NA	NA	NA	NA	NA	NA
EAP.pic	NA	NA	NA	NA	NA	NA	NA	NA	NA
EAP.vr EAP.pic									
EAP.sgp	NA	NA							
EAP.cr	NA	NA							
EAP.dt	NA	NA							
EAP.eg	NA	NA							
EAP.exp	NA	NA							
EAP.ext	NA	NA							
EAP.mod	NA	NA							
EAP.rr	NA	NA							
EAP.tab	NA	NA							
EAP.vr	1	NA							
EAP.pic	NA	1							

Handling missing data

Add `use="complete.obs"` or `use="pairwise.complete.obs"`

```
var(ACEDextract$Correct, ACEDextract$EAP.sgp, use="complete.obs")
```

```
[1] 6.059457
```

```
cor(ACEDextract$Correct, ACEDextract$EAP.sgp, use="complete.obs")
```

```
[1] 0.8359097
```

```
select(ACEDextract, starts_with("EAP")) %>% cor(use="pairwise.complete.obs")
```

	EAP.sgp	EAP.cr	EAP.dt	EAP.eg	EAP.exp	EAP.ext	EAP.mod
EAP.sgp	1.0000000	0.7924686	0.9967152	0.8617104	0.3638555	0.7353437	0.8675896
EAP.cr	0.7924686	1.0000000	0.8150896	0.6401405	0.2575187	0.7654101	0.6497430
EAP.dt	0.9967152	0.8150896	1.0000000	0.8461165	0.3501101	0.7598859	0.8617598
EAP.eg	0.8617104	0.6401405	0.8461165	1.0000000	0.2751036	0.5711900	0.7182038
EAP.exp	0.3638555	0.2575187	0.3501101	0.2751036	1.0000000	0.2124473	0.3428768
EAP.ext	0.7353437	0.7654101	0.7598859	0.5711900	0.2124473	1.0000000	0.6423054
EAP.mod	0.8675896	0.6497430	0.8617598	0.7182038	0.3428768	0.6423054	1.0000000
EAP.rr	0.5163624	0.4335190	0.5178691	0.4484167	0.3499196	0.4315655	0.4758091
EAP.tab	0.7000835	0.5831217	0.7020186	0.5854901	0.3349572	0.6244417	0.5934085
EAP.vr	0.7687073	0.5111927	0.7511615	0.6631146	0.4931869	0.5126654	0.7662889
EAP.pic	0.7055832	0.5067370	0.6867729	0.5971602	0.2377316	0.5008445	0.5979188
	EAP.rr	EAP.tab	EAP.vr	EAP.pic			
EAP.sgp	0.5163624	0.7000835	0.7687073	0.7055832			
EAP.cr	0.4335190	0.5831217	0.5111927	0.5067370			
EAP.dt	0.5178691	0.7020186	0.7511615	0.6867729			
EAP.eg	0.4484167	0.5854901	0.6631146	0.5971602			
EAP.exp	0.3499196	0.3349572	0.4931869	0.2377316			
EAP.ext	0.4315655	0.6244417	0.5126654	0.5008445			
EAP.mod	0.4758091	0.5934085	0.7662889	0.5979188			
EAP.rr	1.0000000	0.4800405	0.4960209	0.3914310			
EAP.tab	0.4800405	1.0000000	0.4766964	0.4328901			
EAP.vr	0.4960209	0.4766964	1.0000000	0.5630053			
EAP.pic	0.3914310	0.4328901	0.5630053	1.0000000			

Biserial and Polyserial correlations

polychor::polychor

Kendal's Tau-b

DescTools::KendalTauB

Fitting a Linear Model

The `lm()` function fits a linear model.

It returns an *object* of class “lm”.

Can do interesting things with the object.

Formulas

The first argument to `lm()` is a formula.

A formula looks like $y \sim x$, where both x and y can be expressions with multiple variables.

`~` is a special character which makes a formula.

y is the dependent variable (what we want to predict)

x is the independent variable (what we are going to use to make the prediction)

Using the example above, `EAP.sgp ~ Correct`.

Generally, it will be the name of a variable, either in the data set or in the global environment.

Can also add a transformation, e.g., `log(x)` or `sqrt(x)`.

Sometimes use a `.` for special purposes.

Other arguments of `lm()`

- **data** – which data set are we using. Name of the data set, or “.” if the data set is being piped in with “%>%”.
- **subset** (optional) – either a vector of cases (row numbers) to use, or a logical vector same as number of rows in data which selects the cases to use.

Also can use `filter()` command on data before `lm()`

- **weights** – normally not used, but support complex survey designs.
- **na.action** – What to do with missing values.
 - “na.fail” – Generate an error
 - “na.omit” – Removes the missing values.
 - “na.exclude” – Removes the missing values, but pads the output so that the missing values can be predicted.
 - “na.pass” – passes the missing values through (result is likely to be NA, so usually not useful).

Can globally set the default by using `options()`

```
options("na.action")
```

```
$na.action  
[1] "na.omit"
```

```
options(na.action=na.fail)
```

```
try(  
lm_EAPxCorrect <- lm(EAP.sgp~Correct, data=ACEDextract)  
)
```

Error in na.fail.default(structure(list(EAP.sgp = c(1.747, 0.000999999999999989, :
missing values in object

```
lm_EAPxCorrect <- lm(EAP.sgp~Correct, data=ACEDextract,  
na.action="na.exclude")
```

Summaries

The result of running `lm` is an S3¹ object of class “lm”.

```
class(lm_EAPxCorrect)
```

```
[1] "lm"
```

Generic functions do things slightly differently based on the, class of the [first] argument.

Methods of S3 generic functions are named `function.class`.

- `print.lm` – `print()` is an important generic function. The `print()` function is called when you just type the name of a variable in the console.

```
lm_EAPxCorrect
```

Call:

```
lm(formula = EAP.sgp ~ Correct, data = ACEDextract, na.action = "na.exclude")
```

Coefficients:

(Intercept)	Correct
-0.67723	0.05446

¹S3 objects, so called because they are described in the 3rd S book, Chambers and Hastie (1992) are lists with a special class attribute.

May want to change the `digits` argument.

```
print(lm_EAPxCorrect,digits=3)
```

Call:

```
lm(formula = EAP.sgp ~ Correct, data = ACEDextract, na.action = "na.exclude")
```

Coefficients:

(Intercept)	Correct
-0.6772	0.0545

- `summary.lm` – The `lm` method of the `summary` function gives the statistics you commonly see in SPSS output.

```
summary(lm_EAPxCorrect)
```

Call:

```
lm(formula = EAP.sgp ~ Correct, data = ACEDextract, na.action = "na.exclude")
```

Residuals:

Min	1Q	Median	3Q	Max
-0.87465	-0.28271	-0.03394	0.23909	1.33593

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.677235	0.058757	-11.53	<2e-16 ***
Correct	0.054456	0.002368	23.00	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.378 on 228 degrees of freedom

(60 observations deleted due to missingness)

Multiple R-squared: 0.6987, Adjusted R-squared: 0.6974

F-statistic: 528.8 on 1 and 228 DF, p-value: < 2.2e-16

- `anova.lm` – This gives the ANOVA table

```
anova(lm_EAPxCorrect)
```

Analysis of Variance Table

Response: EAP.sgp

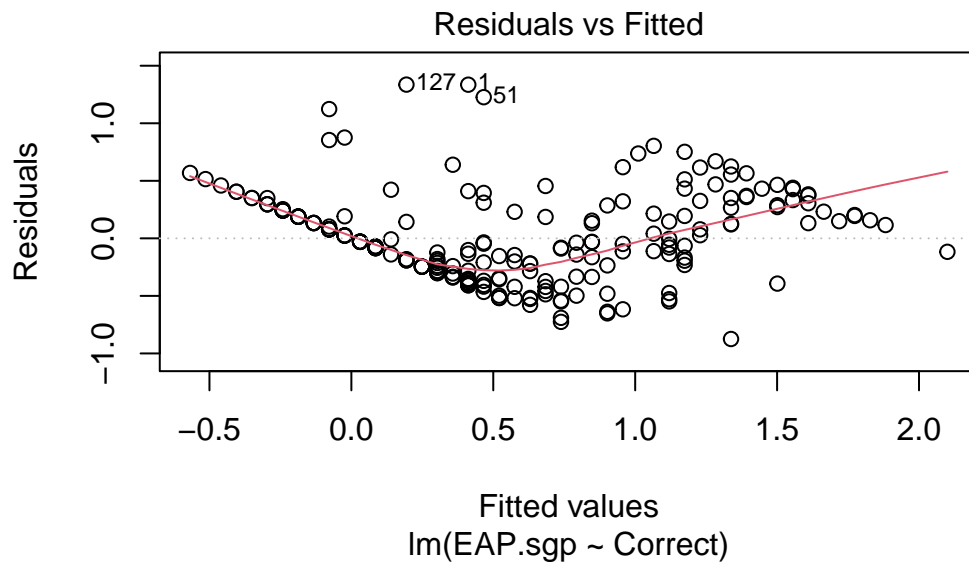
	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Correct	1	75.564	75.564	528.83	< 2.2e-16 ***
Residuals	228	32.579	0.143		

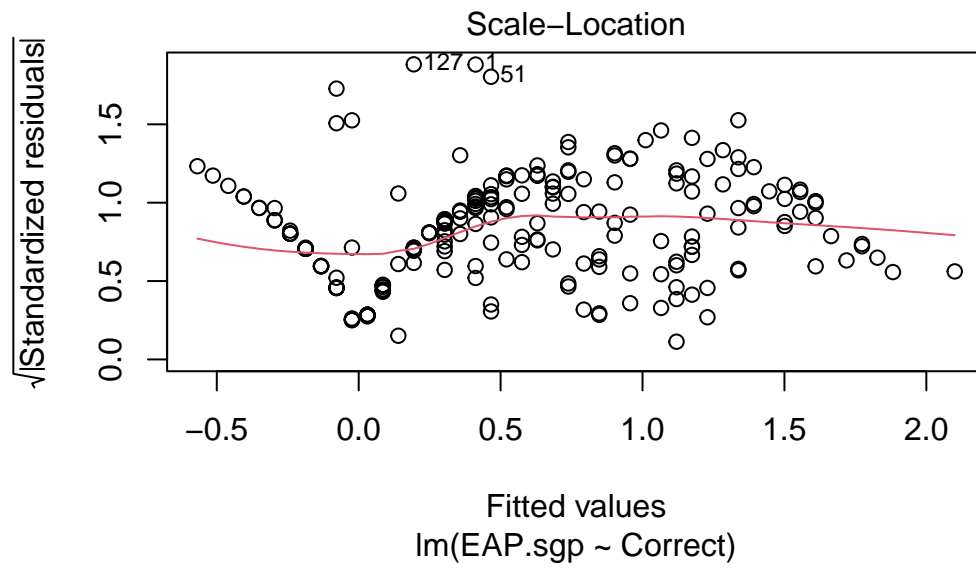

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


- `plot.lm` – This produces a number of diagnostic plots, more later.

By default, the `plot.lm` method asks if you are ready before plotting the next plot. This is not necessary in RStudio, so add the option `ask=FALSE`.

```
plot(lm_EAPxCorrect,ask=FALSE)
```







Note `help(plot)` gives help on the generic (any object) function, and `help(plot.lm)` gives help on the `lm` method for `plot`.

components

An S3 object is basically just a list. To access its components use the `$` operator

- `coefficients` – the slope and intercept
- `residuals` – the vector of residuals
- `fitted.values` – the vector of fitted values
- `df.residuals` – the degrees of freedom of the residuals.

```
lm_EAPxCorrect$coefficients
```

```
(Intercept)      Correct
-0.67723460    0.05445626
```

```
lm_EAPxCorrect$df.residual
```

```
[1] 228
```



```
head(lm_EAPxCorrect$residuals)
```

1	2	3	4	5	6
1.33510938	0.02475946	0.51272172	0.24158451	-0.38589062	0.45482807

```
head(lm_EAPxCorrect$fitted.values)
```

1	2	3	4	5	6
0.41189062	-0.02375946	1.17427828	-0.24158451	0.41189062	0.68417193

- `qr` – The Q and R matrixes from the QR decomposition.

Extracting bits

There are certain common extraction functions. (Usually better to use than the `$` operator.)

- `coef` – coefficients
- `effects` – effects, i.e., coefficients
- `vcov` – variance/covariance matrix
- `nobs` – number of [non-missing] observations.
- `variable.names` – names of variables used in model.

```
coef(lm_EAPxCorrect)
```

(Intercept)	Correct
-0.67723460	0.05445626

```
variable.names(lm_EAPxCorrect)
```

```
[1] "(Intercept)" "Correct"
```

Extracting bits from the summary

- `summary()$sigma` – residual sd/standard error of the estimate
- `summary()$df` – degrees of freedom
- `summary()$fstatistic`
- `summary()$r.squared, summary()$adj.r.squared`

```
summary(lm_EAPxCorrect)$sigma
```

```
[1] 0.3780061
```

```
summary(lm_EAPxCorrect)$r.squared
```

```
[1] 0.6987451
```

Prediction

- `predict`
- `fitted`
- `residuals, rstandard, rstudent`
- `simulate`

Diagnostics

- `dfbeta, dfbetas, dffits`
- `cooks.distance`
- `influence`
- `hatvalues`

Model Fit

- `logLik`
- `deviance`

Tasks

1. Make marginal summaries for the following variables:

Correct,Incorrect,ElapsedTime,pre_scaled,post_scaled,EAP.sgp

```
Desc(ACEDextract$Correct)
```

ACEDextract\$Correct (numeric)


length	n	NAs	unique	0s	mean	meanCI'
	290	230	60	47	0	22.47
		79.3%	20.7%		0.0%	21.10
						23.84
	.05	.10	.25	median	.75	.90
	.95					
	8.00	9.00	14.00	21.00	30.00	37.00
					41.00	
range	sd	vcoef	mad	IQR	skew	kurt
49.00	10.55	0.47	11.86	16.00	0.38	-0.68

lowest : 2.0, 3.0, 4.0, 5.0 (2), 6.0 (2)

highest: 44.0, 45.0 (2), 46.0, 47.0, 51.0

' 95%-CI (classic)


ACEDextract\$Correct (numeric)



2. Same as above, but break down by Level_code
3. Plot Correct against Incorrect. What is happening here?

```
ggplot(ACEDextract,aes(x=Correct,y=Incorrect)) +geom_point(position="jitter")
```

Warning: Removed 60 rows containing missing values (`geom_point()`).



4. Plot EAP.sgp against post_scaled. What is the correlation?
5. Plot pre_scaled against post_scaled. What is the correlation?
6. Regress post_scaled against EAP.sgp. Is EAP.sgp (the internal measure of ability from inside the game) a good predictor of post_scaled (the external predictor)?